

```
# Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Load training and test data
train_data = pd.read_csv("Train.csv")
test_data = pd.read_csv("Test.csv")

# Step 1: Data Preprocessing
# Handle categorical features with Label Encoding
categorical_columns = ['Job', 'Marital', 'Education', 'Default', 'Communication', 'LastContactMonth', 'Outcome']
le = LabelEncoder()

for col in categorical_columns:
    train_data[col] = le.fit_transform(train_data[col].astype(str))
    test_data[col] = le.transform(test_data[col].astype(str))

# Fill missing values for numerical columns, such as Balance, with the median value
train_data['Balance'].fillna(train_data['Balance'].median(), inplace=True)
test_data['Balance'].fillna(test_data['Balance'].median(), inplace=True)

# Step 2: Drop time-based columns (or convert them if needed)
train_data.drop(columns=['CallStart', 'CallEnd'], inplace=True)
test_data.drop(columns=['CallStart', 'CallEnd'], inplace=True)

# Step 3: Feature Scaling
scaler = StandardScaler()
train_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']] = scaler.fit_transform(
    train_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']])
test_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']] = scaler.transform(
    test_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']])

# Step 4: Splitting the training data into train and validation sets
X = train_data.drop('CarInsurance', axis=1)
y = train_data['CarInsurance']

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 5: Bagging Classifier without base_estimator parameter
bagging_model = BaggingClassifier(n_estimators=50, random_state=42)

# Train the model
bagging_model.fit(X_train, y_train)

# Step 6: Evaluate the model
y_pred = bagging_model.predict(X_val)

# Accuracy
accuracy = accuracy_score(y_val, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Classification report
print("Classification Report:")
print(classification_report(y_val, y_pred))

# Confusion Matrix
conf_matrix = confusion_matrix(y_val, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

# Step 7: Feature Importance Visualization from Base Estimators
importances = np.mean([tree.feature_importances_ for tree in bagging_model.estimators_], axis=0)

# Create a DataFrame with feature names and their importances
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importances
})

# Sort the features by importance
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)
```

```
# Plotting feature importances
plt.figure(figsize=(8, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df)
plt.title('Feature Importance')
plt.show()

# Step 8: Predicting on the test data
X_test = test_data.drop('CarInsurance', axis=1)


# Save the 'Id' column separately before dropping it
test_ids = test_data['Id']

# Predict on the test set
y_test_pred = bagging_model.predict(X_test)

# Add predictions to the test data (including the preserved 'Id' column)
test_data['PredictedCarInsurance'] = y_test_pred

# Save predictions along with 'Id' to a CSV file
test_data[['Id', 'PredictedCarInsurance']].to_csv('predictions.csv', index=False)

# Output some predictions (including the 'Id' column)
print(test_data[['Id', 'PredictedCarInsurance']].head())
```

 /var/folders/vg/sx4xytfs3d7259x53nqmn0s80000gn/T/ipykernel_7066/1139527043.py:25: FutureWarning: A value is trying to be set on a mutable object (e.g. a list or ndarray) which is already being modified in place. This behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting the value is a copy of the original object.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
train_data['Balance'].fillna(train_data['Balance'].median(), inplace=True)
```

/var/folders/vg/sx4xytfs3d7259x53nqmn0s80000gn/T/ipykernel_7066/1139527043.py:26: FutureWarning: A value is trying to be set on a mutable object (e.g. a list or ndarray) which is already being modified in place. This behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting the value is a copy of the original object.



For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
test_data['Balance'].fillna(test_data['Balance'].median(), inplace=True)
```

Accuracy: 0.7188

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.85	0.78	484
1	0.69	0.53	0.60	316
accuracy			0.72	800
macro avg	0.71	0.69	0.69	800
weighted avg	0.71	0.72	0.71	800



```
Id PredictedCarInsurance
0 4001 0
1 4002 1
2 4003 1
3 4004 0
4 4005 0
```

```
# Import necessary libraries
import pandas as pd
```

```

import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost as xgb
from sklearn.ensemble import GradientBoostingClassifier

# Load training and test data
train_data = pd.read_csv("Train.csv")
test_data = pd.read_csv("Test.csv")

# Step 1: Data Preprocessing (same as before)
categorical_columns = ['Job', 'Marital', 'Education', 'Default', 'Communication', 'LastContactMonth', 'Outcome']
le = LabelEncoder()

for col in categorical_columns:
    train_data[col] = le.fit_transform(train_data[col].astype(str))
    test_data[col] = le.transform(test_data[col].astype(str))

# Fill missing values for numerical columns, such as Balance, with the median value
train_data['Balance'].fillna(train_data['Balance'].median(), inplace=True)
test_data['Balance'].fillna(test_data['Balance'].median(), inplace=True)

# Step 2: Drop time-based columns (or convert them if needed)
train_data.drop(columns=['CallStart', 'CallEnd'], inplace=True)
test_data.drop(columns=['CallStart', 'CallEnd'], inplace=True)

# Step 3: Feature Scaling
scaler = StandardScaler()
train_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']] = scaler.fit_transform(
    train_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']])
test_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']] = scaler.transform(
    test_data[['Balance', 'Age', 'NoOfContacts', 'DaysPassed', 'PrevAttempts']])

# Step 4: Splitting the training data into train and validation sets
X = train_data.drop('CarInsurance', axis=1)
y = train_data['CarInsurance']

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 5: Train XGBoost Model
xgb_model = xgb.XGBClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)

# Train the XGBoost model
xgb_model.fit(X_train, y_train)

# Step 6: Evaluate XGBoost Model
y_pred_xgb = xgb_model.predict(X_val)

# Accuracy
accuracy_xgb = accuracy_score(y_val, y_pred_xgb)
print(f"XGBoost Accuracy: {accuracy_xgb:.4f}")

# Classification report
print("XGBoost Classification Report:")
print(classification_report(y_val, y_pred_xgb))

# Confusion Matrix for XGBoost
conf_matrix_xgb = confusion_matrix(y_val, y_pred_xgb)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_xgb, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('XGBoost Confusion Matrix')
plt.show()

# Step 7: Train GBM Model
gbm_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)

# Train the GBM model
gbm_model.fit(X_train, y_train)

# Step 8: Evaluate GBM Model
y_pred_gbm = gbm_model.predict(X_val)

# Accuracy
accuracy_gbm = accuracy_score(y_val, y_pred_gbm)
print(f"GBM Accuracy: {accuracy_gbm:.4f}")

# Classification report

```

```

print("GBM Classification Report:")
print(classification_report(y_val, y_pred_gbm))

# Confusion Matrix for GBM
conf_matrix_gbm = confusion_matrix(y_val, y_pred_gbm)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_gbm, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('GBM Confusion Matrix')
plt.show()

# Step 9: Feature Importance Visualization (for XGBoost and GBM)
# XGBoost Feature Importance
xgb_importance = xgb_model.feature_importances_

# GBM Feature Importance
gbm_importance = gbm_model.feature_importances_

# Create DataFrame for XGBoost feature importances
xgb_feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': xgb_importance
})
xgb_feature_importance_df = xgb_feature_importance_df.sort_values(by='Importance', ascending=False)

# Create DataFrame for GBM feature importances
gbm_feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': gbm_importance
})
gbm_feature_importance_df = gbm_feature_importance_df.sort_values(by='Importance', ascending=False)

# Plotting Feature Importances for XGBoost
plt.figure(figsize=(8, 6))
sns.barplot(x='Importance', y='Feature', data=xgb_feature_importance_df)
plt.title('XGBoost Feature Importance')
plt.show()

# Plotting Feature Importances for GBM
plt.figure(figsize=(8, 6))
sns.barplot(x='Importance', y='Feature', data=gbm_feature_importance_df)
plt.title('GBM Feature Importance')
plt.show()

# Step 10: Predicting on the test data (same as before)
X_test = test_data.drop('CarInsurance', axis=1)
test_ids = test_data['Id']

# Predict on the test set using XGBoost
y_test_pred_xgb = xgb_model.predict(X_test)

# Predict on the test set using GBM
y_test_pred_gbm = gbm_model.predict(X_test)

# Add predictions to the test data (including the preserved 'Id' column)
test_data['PredictedCarInsurance_XGB'] = y_test_pred_xgb
test_data['PredictedCarInsurance_GBM'] = y_test_pred_gbm

# Save predictions for XGBoost and GBM to a CSV file
test_data[['Id', 'PredictedCarInsurance_XGB', 'PredictedCarInsurance_GBM']].to_csv('predictions_boosting.csv', index=False)

# Output some predictions (including the 'Id' column)
print(test_data[['Id', 'PredictedCarInsurance_XGB', 'PredictedCarInsurance_GBM']].head())

```



XGBoost Accuracy: 0.7375

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.73	0.91	0.81	484
1	0.77	0.48	0.59	316
accuracy			0.74	800
macro avg	0.75	0.69	0.70	800
weighted avg	0.74	0.74	0.72	800

/var/folders/vg/sx4xytfs3d7259x53nqmn0s80000gn/T/ipykernel_7066/757702331.py:25: FutureWarning: A value is trying to be set on a
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setti

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[


```
train_data['Balance'].fillna(train_data['Balance'].median(), inplace=True)
```

/var/folders/vg/sx4xytfs3d7259x53nqmn0s80000gn/T/ipykernel_7066/757702331.py:26: FutureWarning: A value is trying to be set on a
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setti

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[

```
test_data['Balance'].fillna(test_data['Balance'].median(), inplace=True)
```

XGBoost Confusion Matrix




GBM Accuracy: 0.7375


GBM Classification Report:

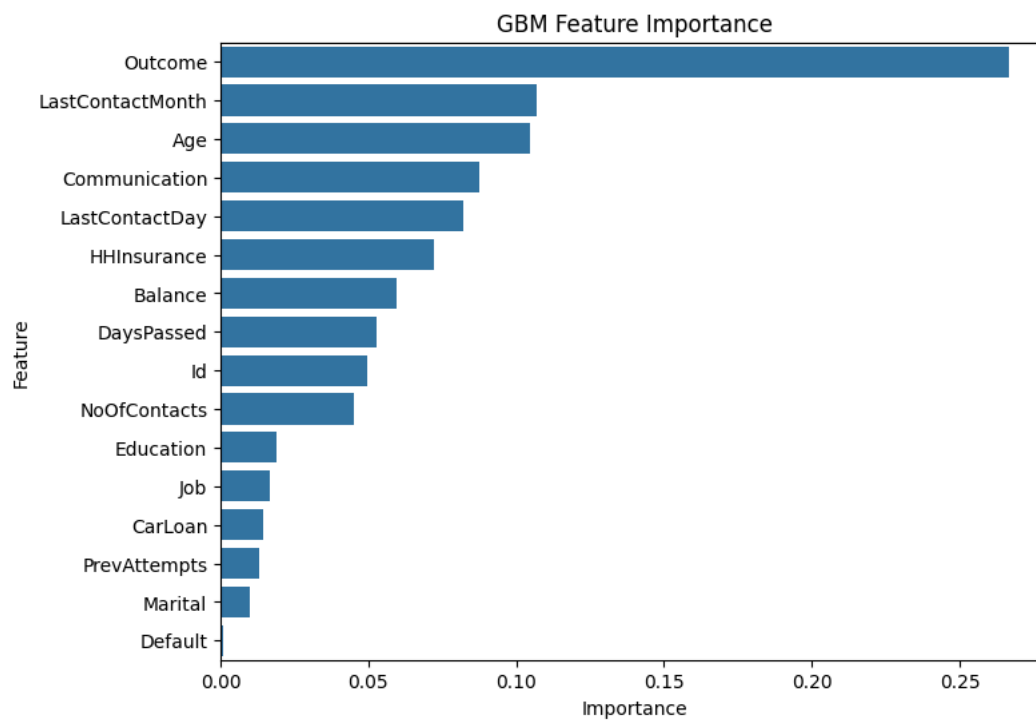

	precision	recall	f1-score	support
0	0.73	0.90	0.81	484
1	0.76	0.49	0.59	316
accuracy			0.74	800
macro avg	0.75	0.69	0.70	800
weighted avg	0.74	0.74	0.72	800

GBM Confusion Matrix



XGBoost Feature Importance





	Id	PredictedCarInsurance_XGB	PredictedCarInsurance_GBM
0	4001	0	0
1	4002	0	0
2	4003	0	0
3	4004	0	0
4	4005	0	0