

# businessproblem

November 20, 2024

```
[1]: # Data Manipulation and Analysis
import pandas as pd
import numpy as np

# Data Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning (Linear Regression and Evaluation)
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[2]: df = pd.read_csv('/Users/vishal/Desktop/CSV files/Labourtraining.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	Age	Education	Race	Hisp	MaritalStatus	Nodeg	\
0	45	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
1	21	Intermediate	NotBlack	NotHispanic	NotMarried	0	
2	38	HighSchool	NotBlack	NotHispanic	Married	0	
3	48	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
4	18	LessThanHighSchool	NotBlack	NotHispanic	Married	1	

	Earnings_1974	Earnings_1975	Earnings_1978
0	21516.670	25243.550	25564.670
1	3175.971	5852.565	13496.080
2	23039.020	25130.760	25564.670
3	24994.370	25243.550	25564.670
4	1669.295	10727.610	9860.869

```
[4]: print(df.columns)
```

```
Index(['Age', 'Education', 'Race', 'Hisp', 'MaritalStatus', 'Nodeg',
      'Earnings_1974', 'Earnings_1975', 'Earnings_1978'],
      dtype='object')
```

```
[5]: df.describe()
```

```
[5]:
```

	Age	Nodeg	Earnings_1974	Earnings_1975	Earnings_1978
count	15992.000000	15992.000000	15992.000000	15992.000000	15992.000000
mean	33.225238	0.295835	14016.800304	13650.803376	14846.659673
std	11.045216	0.456432	9569.795893	9270.403225	9647.391524
min	16.000000	0.000000	0.000000	0.000000	0.000000
25%	24.000000	0.000000	4403.452250	4398.823000	5669.298000
50%	31.000000	0.000000	15123.580000	14557.110000	16421.975000
75%	42.000000	1.000000	23584.180000	22923.737500	25564.670000
max	55.000000	1.000000	25862.320000	25243.550000	25564.670000

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15992 entries, 0 to 15991
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              15992 non-null  int64
1   Education        15992 non-null  object
2   Race             15992 non-null  object
3   Hisp            15992 non-null  object
4   MaritalStatus    15992 non-null  object
5   Nodeg           15992 non-null  int64
6   Earnings_1974    15992 non-null  float64
7   Earnings_1975    15992 non-null  float64
8   Earnings_1978    15992 non-null  float64
dtypes: float64(3), int64(2), object(4)
memory usage: 1.1+ MB
```

### 0.0.1 Handle Missing Values

```
[8]: df = df.dropna()
```

### 0.0.2 Encode Categorical Variables Convert categorical variables (e.g., Race, Hispanic, Married) into numeric form.

```
[10]: # Convert categories into numerical values (e.g., 0 and 1)
df['Race'] = df['Race'].apply(lambda x: 1 if x == 'Black' else 0)
df['Hisp'] = df['Hisp'].apply(lambda x: 1 if x == 'Yes' else 0)
df['MaritalStatus'] = df['MaritalStatus'].apply(lambda x: 1 if x == 'Yes' else 0)
```

```
[11]: print(df[['Race', 'Hisp', 'MaritalStatus']].head())
```

	Race	Hisp	MaritalStatus
0	0	0	0
1	0	0	0

2	0	0	0
3	0	0	0
4	0	0	0

```
[12]: # Features for prediction
X = df[['Age', 'Race', 'Education', 'Hisp', 'MaritalStatus', 'Earnings_1974',
        ↪ 'Earnings_1975']]

# Target variable
y = df['Earnings_1978']
```

```
[13]: print(df.columns)
```

```
Index(['Age', 'Education', 'Race', 'Hisp', 'MaritalStatus', 'Nodeg',
       'Earnings_1974', 'Earnings_1975', 'Earnings_1978'],
      dtype='object')
```

## 1 Train-Test Split

```
[15]: from sklearn.model_selection import train_test_split

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        ↪ random_state=42)
```

## 2 Build and Train the Linear Regression Model

```
[17]: X = df[['Age', 'Race', 'Hisp', 'MaritalStatus', 'Nodeg', 'Earnings_1974',
        ↪ 'Earnings_1975']]
y = df['Earnings_1978']

# Train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        ↪ random_state=42)

# Train the model
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
```

```
[17]: LinearRegression()
```

```
[18]: # Make predictions on the test set
y_pred = model.predict(X_test)
```

### 3 Evaluate the Model

```
[20]: from sklearn.metrics import mean_squared_error, r2_score

# Calculate Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)

# Calculate R-squared Score
r2 = r2_score(y_test, y_pred)
print("R-squared Score:", r2)
```

Mean Squared Error (MSE): 48607834.103273086

R-squared Score: 0.47690656591862246

### 4 Analyze the Coefficients

```
[22]: # Display model coefficients
coefficients = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_
})
print(coefficients)
```

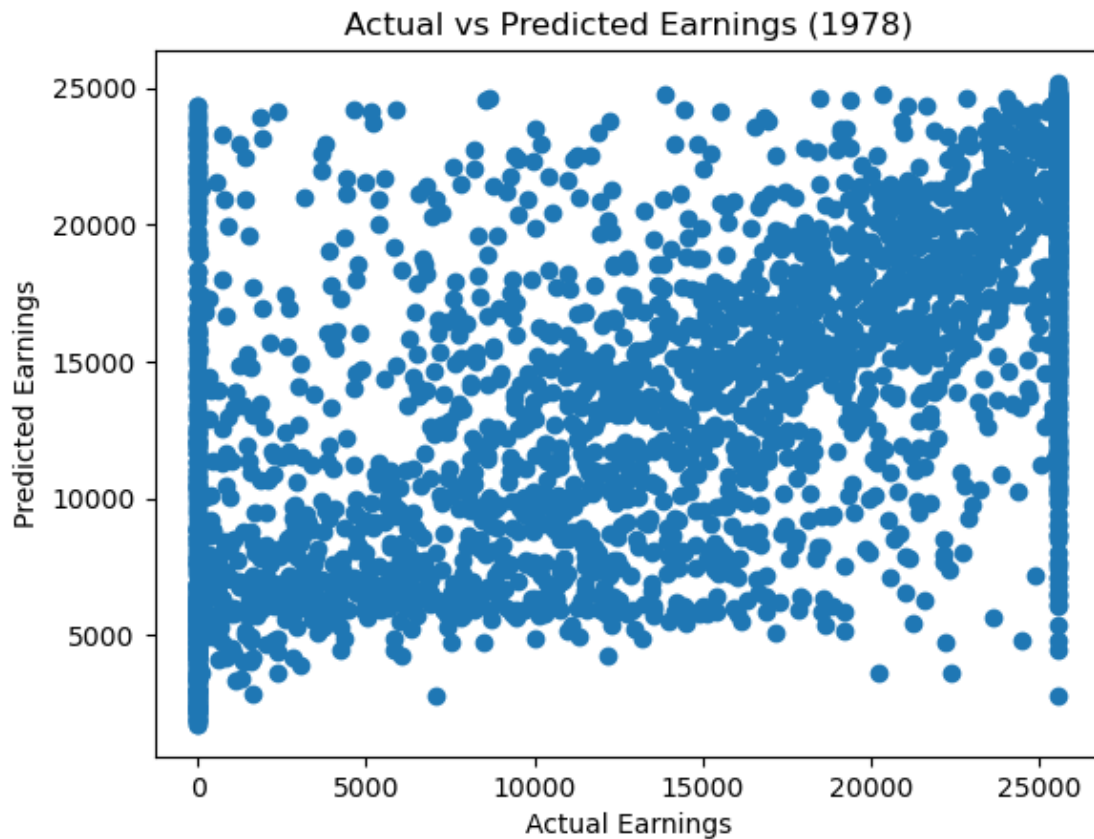
	Feature	Coefficient
0	Age	-1.066479e+02
1	Race	-9.265411e-11
2	Hispanic	-2.246764e-11
3	MaritalStatus	1.705303e-13
4	Nodegree	-4.028111e+02
5	Earnings_1974	2.804689e-01
6	Earnings_1975	4.820691e-01

### 5 Visualize Predictions

### 6 Scatter Plot

```
[25]: import matplotlib.pyplot as plt


plt.scatter(y_test, y_pred)
plt.xlabel("Actual Earnings")
plt.ylabel("Predicted Earnings")
plt.title("Actual vs Predicted Earnings (1978)")
plt.show()
```



## 7 Residual Plot

```
[27]: residuals = y_test - y_pred

plt.scatter(y_pred, residuals)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Predicted Earnings")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```



## 8 Make Predictions on New Data

```
[29]: # Example new data
new_data = pd.DataFrame({
    'Age': [35],
    'Race': [1],
    'Hisp': [0],
    'MaritalStatus': [1],
    'Nodeg': [0],
    'Earnings_1974': [15000],
    'Earnings_1975': [16000]
})

# Predict earnings for 1978
predicted_earnings = model.predict(new_data)
print("Predicted Earnings for 1978:", predicted_earnings)
```

Predicted Earnings for 1978: [16204.28454823]